

TruthLens: A Verification-Gated Pipeline for Evidence-Grounded Fact-Checking of Short-Form Political Video

Aditya Rekhe

[Affiliation placeholder – TODO]

Email: adityarekhe1030@gmail.com

Abstract—Short-form video is now a primary vector for political misinformation, yet most automated fact-checking research targets text – tweets, news sentences, or curated claim lists – and most LLM-based fact-checking systems are evaluated on whether they reach a correct final label, not on whether the intermediate evidence chain that produced that label is real. We present TruthLens, a fact-checking pipeline for short-form political video (Instagram Reels, in the deployed case study) that enforces a strict evidence-before-verdict ordering – atomic claim decomposition, multi-query web research, tiered source retrieval, per-source evidence extraction, and only then a claim-level verdict – and, distinctively, subjects every LLM-generated verdict to a deterministic, non-LLM validation pass before it can be published. The validator checks that every cited piece of evidence actually exists, that every cited source was actually fetched, and that every number asserted in the model’s reasoning actually appears in a cited passage; verdicts that fail are automatically downgraded rather than silently trusted. We further introduce an explicit distinction between UNVERIFIED (evidence was genuinely sought and found insufficient) and RESEARCH_FAILED (the retrieval infrastructure itself broke), so that an infrastructure failure can never be published as if it were a considered null result, and a cost-aware model-escalation pattern in which a free, locally-hosted model is used by default and a paid model is invoked only when a deterministic “is this output substantively grounded” check fails. We report preliminary operational telemetry collected during iterative development – including that deterministic validation downgraded 28.6% (6 of 21) of generated verdicts in our test corpus, and that the free local model required escalation on 25–43% of invocations depending on pipeline stage – and are explicit about what these numbers do and do not establish: they come from a small, non-independent, single-domain test corpus accumulated during debugging, not a held-out benchmark. A small first step toward a real held-out check – two reels TruthLens had never seen, each independently fact-checked by a professional Indian fact-checking organization – found its verdict agreeing with the independent ground truth in both cases, but also surfaced a more consequential result: on both, the system’s text-derived claim extraction missed the actual visual misinformation (real footage recaptioned with a false claim about what it shows), and one run exposed a real gap in validation coverage that let a downgraded verdict’s hallucinated reasoning reach a rendered slide, which we traced and fixed the same day. We describe the larger evaluation this motivates and a taxonomy of concrete failure modes discovered through live testing that we believe generalize to other LLM-pipeline fact-checking systems.

Index Terms—automated fact-checking, large language models, hallucination mitigation, retrieval-augmented generation, misinformation detection, short-form video, claim verification,

source credibility

I. Introduction

Political misinformation increasingly spreads through short-form video – Instagram Reels, YouTube Shorts, TikTok – rather than text [1]. This shift matters for automated fact-checking research in two ways that are easy to underestimate. First, the input is multi-modal and lossy: a claim may be spoken (transcript), displayed as on-screen text (OCR), or only implied by a caption, and any system that treats a reel as a single block of text discards exactly the structure a fact-checker needs to attribute who said what. Second, and more consequentially, the dominant existing evaluation paradigm for automated fact-checking – did the system output the correct veracity label for a claim, measured against a benchmark such as FEVER [2] or LIAR [3] – says nothing about whether the evidence the system cites in support of that label is real. An LLM asked to “fact-check this claim” can produce a fluent, well-cited-looking verdict whose citations point to sources it never retrieved, or whose numbers never appeared in any source it read. This is the specific failure mode this paper is about.

We built and iteratively debugged TruthLens, a pipeline that ingests a political short-form video, decomposes it into atomic factual claims, researches each claim against the open web, and produces a carousel-style, cited, claim-by-claim fact-check – while trying to make that specific failure mode (correct-sounding but ungrounded output) structurally difficult rather than merely discouraged by prompting. The central design principle, stated as a hard constraint throughout development, is evidence before verdict: the system is not permitted to propose a verdict and then search for support; it must search first, and a verdict can only be built from evidence rows that already exist in the database at the time the verdict-proposal model is invoked.

Enforcing this ordering architecturally is necessary but not sufficient, because even a model given only real, already-fetched evidence can still hallucinate – inventing a citation to evidence that was never actually supplied, asserting a number the source never stated, or (a failure mode we discovered live, not anticipated in advance)

padding a reasoning_summary field entirely with internal citation markup and no actual prose, which is schema-valid (a non-empty string) while being substantively empty. TruthLens’s second contribution is therefore a deterministic, non-LLM validation layer that re-checks every LLM output against ground truth already in the database before it is allowed to reach a reader, downgrading rather than trusting outputs that fail.

This paper makes the following contributions:

- 1) An architecture and reference implementation for evidence-grounded, multi-claim fact-checking of short-form video, including atomic claim decomposition with verbatim-quote grounding, multi-tier source credibility scoring, and deterministic (non-LLM) aggregation of per-claim verdicts into a single auditable reel-level verdict (Section III).
- 2) A deterministic post-hoc validation methodology – citation existence, source-fetch existence, and number-grounding checks – that catches a measurable fraction of ungrounded LLM output without requiring a second, equally fallible LLM call to police the first (Section IV).
- 3) An explicit RESEARCH_FAILED vs. UNVERIFIED distinction, and a publication gate that refuses to let infrastructure failure masquerade as a considered null result – an instance of the broader selective-prediction principle that a system should distinguish “I looked and found nothing” from “I could not look” (Section IV).
- 4) A cost-aware escalation pattern – default to a free, local model; escalate to a paid model only when a deterministic substantiveness check fails – evaluated, as preliminary telemetry, on how often escalation actually triggers per pipeline stage (Section V).
- 5) A concrete taxonomy of failure modes discovered through live, iterative testing against a real deployed reel, several of which (a paraphrase misattributed as a direct quote; a dashboard verdict that silently disagreed with the published carousel verdict; a bracket-matching validation-note regex that truncated at the wrong bracket) are not, to our knowledge, previously documented in the automated fact-checking literature at this level of implementation detail (Section VI).

We are explicit throughout about the evidentiary status of our claims. Section V reports real, measured telemetry from TruthLens’s own audit log, not simulated or projected numbers – but that telemetry comes from a small corpus (three distinct source videos, 16 distinct claims, 21 verdict-generation events including re-runs across pipeline-fix iterations) accumulated during debugging rather than a frozen, held-out benchmark, and we describe in Section VIII the larger evaluation required before any accuracy claim here should be treated as representative.

II. Related Work

Automated fact-checking and LLM factuality evaluation have both grown into large enough literatures to warrant their own surveys [4], [5]; we focus below on the specific sub-areas TruthLens draws on and departs from, rather than attempting a comprehensive survey ourselves.

A. Claim detection and decomposition

Claim detection for fact-checking spans monolingual, multilingual, and cross-lingual settings [6], all of which are relevant to TruthLens’s mixed English/Hindi/Urdu input. ClaimBuster [7] established check-worthy claim detection as a distinct pipeline stage, using a TF-IDF/POS/NER-featured SVM to rank sentences from political debate transcripts by how worth fact-checking they are. TruthLens’s claim-extraction stage performs a related but distinct task: rather than ranking existing sentences, an LLM decomposes a transcript, OCR stream, and caption into new atomic claims, splitting compound statements (e.g. “X happened, and because of it Y happened”) into independently checkable parts, following the general claim-decomposition strategy used in systems such as ProgramFC, which decomposes complex claims into sub-questions before per-question retrieval. We extend this with a grounding requirement specific to video: each extracted claim’s source_quote, when present, must be a verbatim substring of the actual transcript/OCR text, checked deterministically, not merely asserted by the model.

B. Fact-checking benchmarks

FEVER [2] (Wikipedia-derived claims, Supports/Refutes/NotEnoughInfo) and LIAR [3] (12.8K PolitiFact statements, six-way truthfulness) are the dominant benchmarks for claim-verification models, and both evaluate a system’s final label against ground truth without examining whether the evidence chain that produced the label is real. TruthLens’s evaluation gap (Section VIII) is precisely that we do not yet have a FEVER/LIAR-scale held-out benchmark for our own domain (Indian political short-form video); closing that gap, likely by curating a small labeled set in the style of these benchmarks rather than reusing them directly (our claims arrive multi-modally and in Hindi/Urdu/English mixed transcripts, which neither benchmark covers), is the single most important item of future work.

C. LLM hallucination and grounding verification

A growing body of work targets citation and grounding hallucination in LLM output specifically. A preregistered, hand-scored evaluation of leading commercial legal AI research tools found each hallucinates on 17–33% of queries [8], and recent systems such as CiteCheck [9] detect hallucinated references by grounding each citation against external sources with a structured LLM verifier. Our validation layer (Section IV) shares the goal but differs in

mechanism: rather than using a second LLM call to judge whether a citation is plausible, TruthLens’s checks are deterministic set-membership and substring tests against data the pipeline itself already fetched (does the cited evidence_id exist; was the cited source actually retrieved; does the asserted number appear in the cited passage’s stored text). This trades generality (we cannot catch a hallucination that happens to reuse a real, cited evidence ID incorrectly) for auditability and zero marginal LLM cost per check. A recent analysis of citation-grounding oracles in the legal domain is a useful caution for exactly this kind of check: it shows that an apparent hallucination rate measured by verifying citations against a reference graph is highly sensitive to that graph’s coverage, not just to the model being measured, with the same responses scoring 15–21% “hallucinated” against a sparse snapshot of a citation registry and under 1.1% against a denser one built from the same source [10]. TruthLens’s own number-grounding check has an analogous blind spot: it can only confirm a number against sources the pipeline itself fetched, so it cannot distinguish “this number is fabricated” from “this number is real but supported only by a source we did not retrieve.” We flag this as an open limitation rather than a solved problem.

D. Political fact-checking with LLMs

Most directly relevant to TruthLens’s domain, DeVerna et al. [11] evaluated 15 LLMs against over 6,000 PolitiFact-verified claims and found that reasoning ability provided “minimal benefit” and raw web search only “moderate gains” for political fact-checking accuracy, while providing curated, high-quality context improved macro-F1 by an average of 233% across models. This is strong independent evidence for TruthLens’s core design bet: that pipeline-level evidence curation (multi-query research planning, source tiering, per-source relevance filtering) matters more than which LLM does the final reasoning step. We view TruthLens as, in part, an applied instantiation and extension of that finding to a harder, multi-modal input domain.

E. Short-form video misinformation

TikTec [1] detects misleading COVID-19 short videos on TikTok by jointly modeling visual, audio, and textual signal, using captions to focus attention on the most informative frames. Recent work on checkworthiness detection specifically targets multilingual short-form video [12]. TruthLens differs from this line of work in task framing: rather than a binary/graded misinformation classifier over the video as a whole, TruthLens performs open-domain, claim-level, cited verification against live web evidence, closer in spirit to FEVER-style verification than to TikTec-style classification, but applied to video-derived claims rather than Wikipedia sentences.

F. Political neutrality of automated fact-checkers

LLMs are not neutral instruments; recent evaluations document measurable partisan asymmetries in frontier models under politically charged prompts [13]. Because a visibly biased automated fact-checker is worse than none for a politically sensitive application, every TruthLens system prompt includes an explicit neutrality clause instructing the model to apply identical evidentiary standards regardless of which political actor is named and to be equally willing to conclude a claim is true or false for any side. We have not run a systematic bias audit of TruthLens’s own output (see Section VIII); this is a limitation, not a resolved claim.

G. Cost-aware model cascading

FrugalGPT [14] demonstrated that routing most queries to a cheap model and escalating only hard cases to an expensive one can sharply cut inference cost with little accuracy loss. TruthLens applies the same principle in reverse-priority form suited to a zero-marginal-cost constraint: every stage defaults to a free, locally hosted model (Llama 3.2 via Ollama) rather than a cheap paid one, and escalates to a paid model (Gemini) only when a deterministic quality gate – not a learned cost/accuracy predictor as in FrugalGPT – detects that the local model’s output is schema-valid but substantively empty.

III. System Architecture

TruthLens is a FastAPI backend with a PostgreSQL/pgvector datastore, a Next.js review dashboard, and an object store (S3-compatible) for media and rendered output. The pipeline runs in the following strict order for every reel, and no stage may begin before the stage before it has persisted its output to the database:

- 1) Ingestion. The source video is downloaded (or manually uploaded), audio is extracted and transcribed, on-screen text is sampled and OCR’d at a fixed interval, and a representative thumbnail frame is selected.
- 2) Claim extraction. An LLM decomposes the transcript, OCR text, and caption into atomic claims, each classified as factual, opinion, prediction, satire, or rhetorical, with only factual claims specific enough to research marked verifiable.
- 3) Research planning. For each verifiable claim, an LLM proposes 2–5 targeted search queries, explicitly instructed to include at least one query aimed at primary sources (official/government domains) when plausible.
- 4) Search and fetch. Each query is run against a search provider; result pages are fetched and their full text extracted (not just the search snippet).
- 5) Evidence analysis. For each fetched source, an LLM judges stance (supports/contradicts/provides_context/irrelevant) against the claim, grounded in that source’s actual extracted text only.

- 6) Verdict proposal. Given the claim and its full evidence matrix (every evidence row, with source reliability and stance), an LLM proposes a verdict label, confidence, and cited evidence IDs.
- 7) Deterministic validation (Section IV).
- 8) Overall-verdict aggregation. Once every claim covered by a fact-check has a validated verdict, a fixed rule table – not an LLM call – derives one reel-level verdict.
- 9) Content assembly and rendering. A carousel of slides (headline poster, claims/quote, evidence, conclusion) is assembled, drawing almost all content directly from already-validated database rows; only headline phrasing and a short “why” paragraph pass through one additional narrow, schema-constrained LLM call each, both with their own grounding checks (Section IV).

A. Source tiering

Every retrieved source is scored along an 8-dimension weighted reliability rubric and assigned to one of eight tiers: `primary_government`, `primary_legal`, `primary_data`, `news_wire`, `established_news`, `academic`, `factcheck_org`, or other. Tier membership determines both which sources are eligible to be surfaced as a claim’s “primary source” citation on the rendered carousel and, in aggregate, the reliability-weighted confidence of the evidence matrix handed to the verdict-proposal stage.

B. Multi-claim, reel-level verdicts

An earlier version of TruthLens built one fact-check per claim; the current version builds one fact-check per reel, covering every verifiable claim extracted from it, because a single short-form video routinely makes several claims of differing veracity, and collapsing them into one per-claim verdict discards exactly the “partially true, partially misleading” structure a reader needs. The rule table in Section IV governs how per-claim verdicts combine into the one verdict actually shown on the published carousel.

IV. Methodology: Making the Model More Trustworthy, Not Just More Capable

The paper’s title question – how to make the model “really better” for this purpose – is, in our experience building TruthLens, substantially a question about the pipeline around the model rather than the model itself. This section describes the four mechanisms we found to matter most.

A. Deterministic anti-hallucination validation

After the verdict-proposal LLM call returns a structured VerdictProposal (label, confidence, reasoning, cited evidence IDs), TruthLens runs three checks against data already in the database, none of which invoke an LLM:

- Citation existence: every ID in `cited_evidence_ids` must correspond to an evidence row that was actually passed to the model in its evidence matrix.
- Source-fetch existence: every cited evidence row’s source must have a stored, successfully fetched full-text passage — a citation cannot point at a source the pipeline only planned to fetch.
- Number grounding: every numeral appearing in the model’s `reasoning_summary` is extracted and checked for appearance in at least one cited source’s stored passage text.

A verdict that fails any check is not discarded – it is downgraded: persisted with a `validation_status` other than passed, an appended, clearly delimited [VALIDATION NOTE: ...] explaining what failed, and excluded from citation lists shown to a reader. This preserves the model’s underlying verdict label (still potentially useful to a human reviewer) while refusing to let its specific, unverified justification reach a reader unflagged.

B. The RESEARCH_FAILED / UNVERIFIED distinction

UNVERIFIED is a legitimate, informative output: it means the system genuinely searched and found insufficient evidence to support or refute a claim. Early in development, TruthLens conflated this with a second, very different situation – the search backend itself returning zero usable results because of a missing API key or provider misconfiguration – and both surfaced identically as UNVERIFIED, indistinguishable to a reader from a considered null result. We introduced a distinct RESEARCH_FAILED claim status, raised only when all queries for a claim fail at the infrastructure level (not when queries succeed but return no relevant results), and a publication gate that refuses to build a fact-check whose primary claim is RESEARCH_FAILED. This is a direct, domain-specific instance of the broader selective-prediction principle that a system should distinguish “I don’t know” from “I could not check”.

C. Verification-gated model cascading

Every pipeline stage defaults to a free, locally hosted model. Three stages – claim extraction, content generation (headline/why-paragraph), and verdict proposal – additionally run a deterministic substantiveness check on the local model’s output before accepting it:

- Claim extraction: is the extraction grounded, i.e. do at least half of the extracted verifiable claims have a `source_quote` that is an actual substring of the combined transcript/OCR/caption text?
- Content generation / verdict reasoning: after stripping internal citation markup ([`evidence_id=...`] blocks), does the remaining prose contain at least six real words? This check exists because we discovered, live, that a locally hosted model would sometimes produce a `reasoning_summary` composed entirely of citation markup with no prose at all – schema-valid (a non-empty string), but explaining nothing to a reader.

On failure, the same call is retried once against a stronger paid model (Gemini) with an identical prompt, and the retry (whichever model produced it) is logged. This differs from a classic cascade [14] in that the escalation trigger is a deterministic grounding/substantiveness predicate specific to this domain’s failure modes, not a learned or generic confidence estimate.

D. Deterministic overall-verdict aggregation

Once every claim covered by a fact-check has a verdict, TruthLens derives one reel-level verdict via a fixed rule table over the set of per-claim verdict labels (weighted by whether a claim is “major,” importance ≥ 0.5) – for example, all claims TRUE yields overall TRUE; a mix containing both TRUE/MOSTLY_TRUE and FALSE yields MISLEADING. This is a plain combinatorial function, not a further LLM call, so the single highest-stakes number in the product – the verdict shown on the published carousel’s first slide – remains fully auditable from already-validated per-claim verdicts rather than resting on a second, opaque judgment.

V. Evaluation

A. What this section is, and is not

We report real telemetry recorded in TruthLens’s own audit log during iterative development and live debugging against a small corpus: five ingested reel records over three distinct underlying videos (by content hash), yielding 28 extracted claims (16 marked verifiable and carried through to a verdict at least once), and 21 verdict-generation events in total – more than the number of distinct claims because several claims were re-verdicted after pipeline fixes across iterations. This is operational telemetry from a system under active development, not a held-out benchmark evaluation: the sample is small, the units are not independent (the same claim appears multiple times across pipeline-fix iterations), and every video in the corpus is Indian political content in a mix of English, Hindi, and Urdu. We report it because the numbers are real and, we believe, still informative about where a deterministic-validation, verification-gated-cascade approach actually spends its safety margin – but no accuracy or rate reported here should be read as representative of held-out performance. Table I summarizes the raw counts referenced throughout this section; all figures are computed directly from the verdicts, claims, evidence, sources, and audit_logs tables described in Section III.

B. Deterministic validation catch rate

Of 21 verdict-generation events, 15 (71.4%) passed all three deterministic checks unmodified; 6 (28.6%) were downgraded – 5 for an unsupported statistic (a number in the reasoning not found in any cited passage) and 1 for a missing/invalid citation. At the distinct-claim level, 5 of 16 claims (31.3%) had at least one verdict downgraded during testing. We read this as evidence that the validation layer

TABLE I
Corpus and telemetry summary (development-time, not a held-out benchmark)

Distinct source videos (by content hash)	3
Reel ingestion records	5
Extracted claims	28
Distinct claims with ≥ 1 verdict	16
Verdict-generation events (incl. re-runs)	21
Evidence rows analyzed	72
– classified supports	22
– classified contradicts	10
– classified irrelevant	40
Logged search-query failures	55
Logged successful evidence-fetch completions	20
Total tokens logged (input / output)	123,477 / 19,821

described in Section IV-A is not a redundant formality: on this corpus, it materially changed what was published for roughly three in ten verdict-generation events, every one of which would otherwise have been schema-valid and passed through unflagged.

C. Verification-gated cascade trigger rate

The local model’s output failed the deterministic substantiveness check, triggering escalation to Gemini, in 3 of 7 logged claim-extraction completions (approx. 43%), 2 of 5 content-generation completions (40%), and 2 of 8 verdict-proposal completions (25%). These are small counts and should not be read as a precise escalation rate for the underlying local model in general; they are, however, real, and large enough to show the gate firing routinely rather than rarely – a free-only, no-escalation deployment on this corpus would have shipped several outputs that were schema-valid but substantively empty.

D. Evidence signal-to-noise

Of 72 evidence rows the evidence-analysis stage judged, 40 (55.6%) were classified irrelevant – the fetched source, despite being returned by a search query targeted at the specific claim, did not actually address it. This is consistent with our qualitative finding (Section VI) that search recall, not claim reasoning, is the current bottleneck: over half of retrieved-and-read pages contribute nothing to the eventual verdict.

E. Source tier distribution

Of sources actually retrieved and scored, the majority (54 of 72 scored sources) fell into the lowest other tier (mean reliability 0.404), with only 8 reaching primary_government (mean reliability 0.674) and none reaching primary_legal or primary_data in this corpus, despite the research-planning prompt explicitly requesting at least one primary-source-targeted query per claim. This gap – the prompt asks for primary sources, but the free search backend rarely surfaces them, especially for India-specific government domains – is, in our assessment, the most consequential open problem for output quality,

ahead of anything in the LLM reasoning stages themselves; we return to it in Section VIII.

F. An out-of-sample pilot

Every number above comes from the same handful of videos used during development. As a first, small step toward the held-out evaluation Section VIII calls for, we ran TruthLens against two reels it had never seen, each independently fact-checked by a professional Indian fact-checking organization (BOOM Live, Alt News), sourced by fetching the organization’s own article and extracting the original Instagram post it had embedded as evidence – not by us selecting or writing the claim or the verdict. On both, TruthLens’s reel-level verdict (MOSTLY_FALSE in both cases) agreed with the independent ground truth (FALSE in both cases). We do not read this as a validated accuracy result – $n = 2$ – but the process of running it surfaced something more informative than the label match itself.

On the second case, the label match was coincidental rather than earned. The reel’s caption never states the claim that was actually false in words; the misinformation is entirely visual (real World Cup celebration footage captioned as a protest). TruthLens’s text-derived claim extraction found two tangential claims in the caption’s moralizing language and never identified or checked the one thing that was actually misleading: what the video shows. Both pilot cases follow this same pattern – genuine footage, false caption about what it depicts – which was also the dominant pattern among the professional fact-checks we surveyed while sourcing this pilot. TruthLens has no video-provenance or reverse-image-search capability; its architecture may be structurally weak against the single most common real-world failure mode in this domain, independent of anything about LLM reasoning quality. We treat this as the most important finding of this pilot and return to it in Section VIII.

The same run also surfaced two further bugs, both fixed the same day (Section VI): evidence retrieval confused two similarly-named but distinct organizations, and a verdict the deterministic validator had already downgraded still had its full free-text reasoning reused downstream, carrying an entity-confused, unsupported claim about a real organization’s ideology onto a rendered slide. Neither is visible in the aggregate statistics above, which is itself worth noting: a system can look healthy on its summary numbers while a specific, reproducible, reputationally serious failure sits one level down.

VI. A Taxonomy of Failure Modes Found in Live Testing

Beyond the aggregate statistics above, iterative live testing against a real deployed reel surfaced several specific, reproducible failure modes that we believe are under-documented in the automated fact-checking literature at this level of implementation detail, because they only appear when a system is actually run end-to-end against real content rather than evaluated on a static benchmark.

- Schema-valid, substantively empty output. Discovered independently in three different pipeline stages (claim extraction, content generation, verdict reasoning): a locally hosted model produces output that satisfies its Pydantic/JSON schema – a non-empty string, a well-typed list – while containing no actual information, e.g. a reasoning_summary composed entirely of `[[evidence_id=...]]` citation markup with zero prose. No schema validation catches this class of failure; only a content-level substantiveness check does (Section IV-C).
- Verbatim-but-misattributed quotation. A claim’s source_quote can be a genuine, verbatim substring of the video’s on-screen text – passing a naive “is this a real quote” check – while still being the wrong kind of quote: a news-style caption written by the video’s publisher about a person (“X Attacks Y Over Z Rule”), not words that person said. Attributing such a quote to the person it names, because their name appears nearby, produces a fabricated first-person quotation from genuine, verbatim text. We fixed this by requiring the candidate speaker’s own name not to appear inside the quote text itself before trusting the attribution – a person’s own quote essentially never refers to them self by full name.
- Silent verdict/presentation inconsistency. After introducing reel-level, multi-claim verdict aggregation (Section III), the review dashboard continued displaying the primary claim’s own individual verdict as “the” verdict, while the actually-published carousel displayed the separately computed reel-level verdict. The two could and did disagree (one test case: dashboard showed the primary claim’s UN-VERIFIED; the carousel showed the reel’s aggregate MOSTLY_TRUE) with no error, because both values were valid data – the bug was in which one the interface trusted, not in either computation. This class of failure – two correct numbers, wrong place shown – will not be caught by any per-field validation test and was only found by comparing rendered output against the database directly.
- Bracket-matching truncation bug. A display-sanitization regex intended to strip an internal [VALIDATION NOTE: ...] suffix used non-greedy bracket matching, which stopped at the first] character encountered – the closing bracket of a Python list repr inside the note itself (e.g. “Numbers [‘3065258’, ‘69’] in reasoning_summary...”), leaving the note’s own tail visible on a published slide. Fixed by truncating at a literal marker string rather than attempting to bracket-match free-form text.
- Unrepresentative thumbnail selection. A single frame extracted at a hardcoded 0.5-second offset is not a reliable representative thumbnail; in one case it landed on an unflattering extreme close-up mid-gesture. We now sample several candidate frames

across the video’s duration and select by a cheap grayscale-variance “visual complexity” score, screening out blank/flat/transition frames – not real scene- or face-aware selection, but a materially better floor than one fixed guess.

- Entity confusion between similarly-named organizations. Found during the out-of-sample pilot (Section V-F): researching a claim about “Karni Sena,” evidence retrieval returned a Wikipedia article about Sri Ram Sena – a real but different Indian organization sharing the common suffix “Sena.” The evidence-analysis stage labeled this source’s stance as contradicts rather than irrelevant: an article about a different entity does not contradict a claim about this one, it simply does not address it. Name-collision entity disambiguation is not currently checked anywhere in the pipeline.
- Downgraded reasoning reused as trusted input downstream. Also found during the out-of-sample pilot, and the more serious of the two: the verdict-proposal model, given the mismatched Sri Ram Sena evidence above, wrote that this “suggests an unrelated organization with violent ideology” – a real hallucination the deterministic validator correctly downgraded (for an unrelated reason: an unsupported number elsewhere in the same reasoning text). But downgrading a verdict only forces its label to UNVERIFIED; it does not redact the free-text reasoning_summary the model wrote to justify its original, rejected verdict, because there is no reliable way to isolate just the offending sentence from otherwise -reasonable-looking prose. That full text was still being reused verbatim in two places – an evidence card’s displayed answer text, and as raw input material for the separate overall “why” paragraph LLM call, which wove it into the actual rendered slide. This is a real instance of exactly the failure class the validation architecture (Section IV-A) exists to prevent, in a genuine gap in that architecture’s coverage: validation applied to whether a verdict is safe to label, not to whether its reasoning text is safe to quote. Fixed by treating any non-passed validation_status as disqualifying for reuse of that verdict’s free text anywhere, replacing it with a generic, non-fabricated note instead. This never reached a real audience – the fact-check stayed in the internal review queue – but the underlying entity-confusion in evidence retrieval is unfixed and could resurface in a form these particular checks don’t catch.

We consider this taxonomy, and the general lesson that most of these failures were only discoverable by running the full system against real content and inspecting rendered output rather than by unit-testing individual functions in isolation, to be as practically important a contribution of this work as any single architectural

component.

VII. Discussion and Threats to Validity

Small, non-independent sample. The telemetry in Section V is drawn from repeated testing against a handful of real videos, not independent draws from a target population; several statistics (e.g. the 28.6% validation-downgrade rate) could shift substantially with a larger, held-out corpus. We report exact counts, not just percentages, throughout so a reader can judge this directly.

Single domain, single language mix. All evaluated content is Indian political short-form video with English/Hindi/Urdu mixed transcripts. Source-tiering domain lists and the research-planning prompt’s primary-source guidance were not tuned for, and have not been tested against, other national contexts.

No systematic bias audit. Section II discusses documented partisan asymmetry in frontier LLMs; TruthLens mitigates this only via prompt-level instruction (the neutrality clause), which is not verified against a labeled bias-evaluation set.

No comparison baseline. We have not run a naive single-shot “ask an LLM to fact-check this claim” baseline under identical conditions to quantify how much the pipeline architecture (versus the underlying model) contributes to output quality. This is the most direct way to substantiate the paper’s central claim and is the top item in Section VIII.

VIII. Future Work

- 1) A held-out, labeled benchmark. Curate a labeled set of short-form political video claims, disjoint from any content used during development, with ground-truth verdicts from independent human reviewers (ideally with inter-annotator agreement reported), and report precision/recall/F1 against it – the single change that would most improve this paper’s evidentiary standing.
- 2) A naive-baseline comparison. Run a single-shot, no-pipeline LLM fact-check over the same claims and report the validation-downgrade rate and accuracy difference, to isolate the contribution of pipeline architecture from the contribution of the underlying model.
- 3) Video-provenance verification. The out-of-sample pilot (Section V-F) found both tested cases, and the majority of professional fact-checks surveyed while sourcing it, follow the same pattern: genuine footage recaptioned with a false claim about what it shows or where it’s from. TruthLens has no mechanism to check this – claim extraction works from transcript/OCR/caption text, and a purely-visual false claim (wrong location, wrong event, wrong date) is invisible to it regardless of how good the downstream reasoning is. Reverse-image/video search against known reference footage, or at minimum flagging

when a claim about a video’s provenance cannot be checked from text alone (rather than silently checking a different, tangential claim instead), would target what the pilot suggests is the dominant real-world failure mode in this domain – likely higher leverage than any of the LLM-reasoning-focused items below.

- 4) Primary-source retrieval. Section V’s most actionable finding: only 8 of 72 retrieved sources in our corpus reached the `primary_government` tier and none reached `primary_legal/primary_data`, despite explicit prompt-level requests for primary-source queries. Concrete next steps: structured queries against known government/legislative document APIs and domain allow-lists (e.g. `.gov.in`, `pib.gov.in`, `indiankanoon.org`) rather than relying on a general-purpose web search engine to surface them incidentally.
- 5) Systematic bias evaluation. Apply an existing partisan- asymmetry evaluation methodology (Section II) to TruthLens’s own verdicts across politically matched claim pairs.
- 6) Cross-lingual and cross-domain generalization. Test source tiering and research-planning prompts against non-Indian political content and languages beyond the current English/Hindi/Urdu mix.
- 7) Calibration. Compare TruthLens’s reported confidence bands against empirical correctness rates once a labeled benchmark exists (Item 1), to check whether “very high confidence” verdicts are in fact more often correct than “moderate confidence” ones.

IX. Conclusion

TruthLens demonstrates that a substantial fraction of the “does the model hallucinate” problem in LLM-based fact-checking is addressable architecturally rather than only through better prompting or a more capable model: a deterministic, non-LLM validation layer applied after generation caught 28.6% of generated verdicts in our development corpus, and an explicit distinction between genuine null results and infrastructure failure prevents an entire class of false confidence. These results come from a small, non-independent development corpus and should be read as preliminary telemetry, not a benchmark evaluation – but the failure modes they surfaced (schema-valid-but-empty output, verbatim-but-misattributed quotation, silent presentation inconsistency) are, we believe, general enough to matter to any team building an LLM pipeline whose output will be published and trusted. The most direct path to strengthening every claim in this paper is the held-out benchmark and naive-baseline comparison described in Section VIII, which we intend to pursue as this work continues.

References

- [1] L. Shang, Z. Kou, Y. Zhang, and D. Wang, “A multimodal misinformation detector for COVID-19 short videos on TikTok,”

- in 2021 IEEE International Conference on Big Data (Big Data), 2021.
- [2] J. Thorne, A. Vlachos, C. Christodoulopoulos, and A. Mittal, “FEVER: a large-scale dataset for fact extraction and VERification,” in Proceedings of the 2018 Conference of the North American Chapter of the Association for Computational Linguistics: Human Language Technologies (NAACL-HLT), 2018.
- [3] W. Y. Wang, “‘Liar, Liar Pants on Fire’: A new benchmark dataset for fake news detection,” in Proceedings of the 55th Annual Meeting of the Association for Computational Linguistics (Volume 2: Short Papers), Vancouver, Canada, 2017.
- [4] S. S. Rahman, M. A. Islam, M. M. Alam, M. Zeba, M. A. Rahman, S. S. Chowdhury, M. A. K. Raiaan, and S. Azam, “Hallucination to truth: A review of fact-checking and factuality evaluation in large language models,” *Artificial Intelligence Review* (Springer); arXiv preprint arXiv:2508.03860, 2025.
- [5] I. Vykopal, M. Pikuliak, S. Ostermann, and M. Šimko, “Generative large language models in automated fact-checking: A survey,” arXiv preprint arXiv:2407.02351, 2024.
- [6] R. Panchendrarajan and A. Zubiaga, “Claim detection for automated fact-checking: A survey on monolingual, multilingual and cross-lingual research,” arXiv preprint arXiv:2401.11969, 2024.
- [7] N. Hassan, F. Arslan, C. Li, and M. Tremayne, “Toward automated fact-checking: Detecting check-worthy factual claims by ClaimBuster,” in Proceedings of the 23rd ACM SIGKDD International Conference on Knowledge Discovery and Data Mining (KDD), 2017, pp. 1803–1812.
- [8] V. Magesh, F. Surani, M. Dahl, M. Suzgun, C. D. Manning, and D. E. Ho, “Hallucination-free? assessing the reliability of leading AI legal research tools,” arXiv preprint arXiv:2405.20362, 2024, preregistered, hand-scored (202 queries) evaluation finding leading commercial legal AI tools hallucinate 17–33% of the time.
- [9] K. Khajavi, S. Sadeghi, R. Adhikari, and A. Tessier, “CiteCheck: Retrieval-grounded detection of LLM citation hallucinations in scientific text,” arXiv preprint arXiv:2605.27700, 2026.
- [10] V. Ovcharov, “Citation grounding measures the oracle: Graph coverage determines reported LLM hallucination rates in law,” arXiv preprint arXiv:2606.00898, 2026, shows apparent citation-hallucination rates from graph-based oracles (15–21% on a sparse snapshot) are largely an artifact of citation-graph coverage rather than a stable property of the model being measured.
- [11] M. R. DeVerna, K.-C. Yang, H. Y. Yan, and F. Menczer, “Large language models require curated context for reliable political fact-checking – even with reasoning and web search,” arXiv preprint arXiv:2511.18749, 2025.
- [12] H. Vatndal and V. Setty, “ShortCheck: Checkworthiness detection of multilingual short-form videos,” in Proceedings of the 14th International Joint Conference on Natural Language Processing and the 4th Conference of the Asia-Pacific Chapter of the Association for Computational Linguistics: System Demonstrations, Mumbai, India, 2025.
- [13] T.-Q. Peng, K. Yang, S. Lee, H. Li, Y. Chu, Y. Lin, and H. Liu, “Beyond partisan leaning: A comparative analysis of political bias in large language models,” *Journal of Information Technology & Politics*, 2026.
- [14] L. Chen, M. Zaharia, and J. Zou, “FrugalGPT: How to use large language models while reducing cost and improving performance,” arXiv preprint arXiv:2305.05176, 2023.